
Problem Set 3: Game of Fifteen

This is CS50. Harvard University.

Table of Contents

| | |
|----------------------------------|----|
| Objectives | 1 |
| Recommended Reading | 1 |
| diff pset3.pdf hacker3.pdf | 2 |
| Academic Honesty | 2 |
| Reasonable | 3 |
| Not Reasonable | 3 |
| Getting Ready | 4 |
| Getting Started | 5 |
| Searching | 7 |
| search | 12 |
| Sorting | 14 |
| sort | 14 |
| search | 16 |
| The Game Begins | 16 |

This is the Hacker Edition of Problem Set 1. It cannot be submitted for credit.

Objectives

- Accustom you to reading someone else's code.
- Introduce you to larger programs and programs with multiple source files.
- Empower you with Makefiles.
- Implement a party favor.

Recommended Reading

- Page 17 of <http://www.howstuffworks.com/c.htm>.

- Chapters 20 and 23 of *Absolute Beginner's Guide to C*.
- Chapters 13, 15, and 18 of *Programming in C*.

[diff pset3.pdf](#) [hacker3.pdf](#)

- Hacker Edition dares you to implement sort in $O(n)$ instead of $O(n^2)$.
- Hacker Edition asks you to play God.

Academic Honesty

This course's philosophy on academic honesty is best stated as "be reasonable." The course recognizes that interactions with classmates and others can facilitate mastery of the course's material. However, there remains a line between enlisting the help of another and submitting the work of another. This policy characterizes both sides of that line.

The essence of all work that you submit to this course must be your own. Collaboration on problem sets is not permitted except to the extent that you may ask classmates and others for help so long as that help does not reduce to another doing your work for you. Generally speaking, when asking for help, you may show your code to others, but you may not view theirs, so long as you and they respect this policy's other constraints. Collaboration on quizzes is not permitted at all. Collaboration on the course's final project is permitted to the extent prescribed by its specification.

Below are rules of thumb that (inexhaustively) characterize acts that the course considers reasonable and not reasonable. If in doubt as to whether some act is reasonable, do not commit it until you solicit and receive approval in writing from the course's heads. Acts considered not reasonable by the course are handled harshly. If the course refers some matter for disciplinary action and the outcome is punitive, the course reserves the right to impose local sanctions on top of that outcome that may include an unsatisfactory or failing grade for work submitted or for the course itself.

If you commit some act that is not reasonable but bring it to the attention of the course's heads within 72 hours, the course may impose local sanctions that may include an unsatisfactory or failing grade for work submitted, but the course will not refer the matter for further disciplinary action except in cases of repeated acts.

Reasonable

- Communicating with classmates about problem sets' problems in English (or some other spoken language).
- Discussing the course's material with others in order to understand it better.
- Helping a classmate identify a bug in his or her code at office hours, elsewhere, or even online, as by viewing, compiling, or running his or her code, even on your own computer.
- Incorporating snippets of code that you find online or elsewhere into your own code, provided that those snippets are not themselves solutions to assigned problems and that you cite the snippets' origins.
- Reviewing past semesters' quizzes and solutions thereto.
- Sending or showing code that you've written to someone, possibly a classmate, so that he or she might help you identify and fix a bug.
- Sharing snippets of your own code online so that others might help you identify and fix a bug.
- Turning to the web or elsewhere for instruction beyond the course's own, for references, and for solutions to technical difficulties, but not for outright solutions to problem set's problems or your own final project.
- Whiteboarding solutions to problem sets with others using diagrams or pseudocode but not actual code.
- Working with (and even paying) a tutor to help you with the course, provided the tutor does not do your work for you.

Not Reasonable

- Accessing a solution to some problem prior to (re-)submitting your own.
- Asking a classmate to see his or her solution to a problem set's problem before (re-)submitting your own.
- Decompiling, deobfuscating, or disassembling the staff's solutions to problem sets.
- Failing to cite (as with comments) the origins of code or techniques that you discover outside of the course's own lessons and integrate into your own work, even while respecting this policy's other constraints.

Problem Set 3:
Game of Fifteen

- Giving or showing to a classmate a solution to a problem set's problem when it is he or she, and not you, who is struggling to solve it.
- Looking at another individual's work during a quiz.
- Paying or offering to pay an individual for work that you may submit as (part of) your own.
- Providing or making available solutions to problem sets to individuals who might take this course in the future.
- Searching for, soliciting, or viewing a quiz's questions or answers prior to taking the quiz.
- Searching for or soliciting outright solutions to problem sets online or elsewhere.
- Splitting a problem set's workload with another individual and combining your work.
- Submitting (after possibly modifying) the work of another individual beyond allowed snippets.
- Submitting the same or similar work to this course that you have submitted or will submit to another.
- Submitting work to this course that you intend to use outside of the course (e.g., for a job) without prior approval from the course's heads.
- Using resources during a quiz beyond those explicitly allowed in the quiz's instructions.
- Viewing another's solution to a problem set's problem and basing your own solution on it.

Getting Ready

First, re-acquaint yourself with our old friends, linear search and binary search (and Patrick), if hazy:

<https://www.youtube.com/watch?v=CX2CYIJLwfg>

Next, get to know our new friends, bubble sort, selection sort, insertion sort, and merge sort (and Jackson, Tommy, and Rob):

<https://www.youtube.com/watch?v=8Kp-8OGwphY>

Phew, so many shorts. And so many sorts! Ha.

Problem Set 3: Game of Fifteen

Finally, a more in-depth look at debugging techniques from Dan! (Odds are these 23 minutes with Dan will save you hours over the course of the term, since GDB is a far better tool than `printf` in many cases!)

https://www.youtube.com/watch?v=-G_kIBQLgdc

Be sure you're reasonably comfortable answering the below when it comes time to submit this problem set's form!

- Why does binary search require that an array be sorted?
- Why is bubble sort in $O(n^2)$?
- Why is insertion sort in $\Omega(n)$?
- In no more than 3 sentences, how does selection sort work?
- What's an upper bound on the (worst-case) running time of merge sort?
- GDB lets you "debug" program, but, more specifically, what does it let you do?

Getting Started

Recall that, for Problem Sets 1 and 2, you started writing programs from scratch, creating your own `pset1` and `pset2` directories with `mkdir`. (And for Problem Set 0, you started writing programs in Scratch! Ha.) For Problem Set 3, you'll instead download "distribution code" (otherwise known as a "distro"), written by us, and add your own lines of code to it. You'll first need to read and understand our code, though, so this problem set is as much about learning to read someone else's code as it is about writing your own!

Let's get you started. Log into cs50.io¹ and execute

```
update50
```

within a terminal window to make sure your workspace is up-to-date. If you somehow closed your terminal window (and can't find it!), make sure that **Console** is checked under the **View** menu, then click the green, circled plus (+) in CS50 IDE's bottom half, then select **New Terminal**.

Next, execute

¹ <https://cs50.io/>

Problem Set 3: Game of Fifteen

`cd ~/workspace`

at your prompt to ensure that you're inside of `workspace` (which is inside of your home directory). Then execute

`wget http://cdn.cs50.net/2015/fall/psets/3/hacker3/hacker3.zip`

to download a ZIP of this problem set's distro into your workspace (with a command-line program called `wget`). You should see a bunch of output followed by:

`'hacker3.zip' saved`

Confirm that you've indeed downloaded `hacker3.zip` by executing

`ls`

and then run

`unzip hacker3.zip`

to unzip the file. If you then run `ls` again, you should see that you have a newly unzipped directory called `hacker3` as well. Proceed to execute

`cd hacker3`

followed by

`ls`

and you should see that the directory contains two "subdirectories":

`fifteen find`

How fun!

Searching

Okay, let's dive into one of those subdirectories. Execute the command below in a terminal window.

```
cd ~/workspace/hacker3/find
```

If you list the contents of this directory (remember how?), you should see the below.

```
helpers.c  helpers.h  Makefile  find.c  generate.c
```

Wow, that's a lot of files, eh? Not to worry, we'll walk you through them.

Implemented in `generate.c` is a program that uses a "pseudorandom-number generator" (via a function called `drand48`) to generate a whole bunch of random (well, pseudorandom, since computers can't actually generate truly random) numbers, one per line. Go ahead and compile this program by executing the command below.

```
make generate
```

Now run the program you just compiled by executing the command below.

```
./generate
```

You should be informed of the program's proper usage, per the below.

```
Usage: generate n [s]
```

As this output suggests, this program expects one or two command-line arguments. The first, `n`, is required; it indicates how many pseudorandom numbers you'd like to generate. The second, `s`, is optional, as the brackets are meant to imply; if supplied, it represents the value that the pseudorandom-number generator should use as its "seed." A seed is simply an input to a pseudorandom-number generator that influences its outputs. For instance, if you seed `drand48` by first calling `srand48` (another function whose purpose is to "seed" `drand48`) with an argument of, say, `1`, and then call `drand48` itself three times, `drand48` might return `2728`, then `29785`, then `54710`. But if you instead seed

Problem Set 3: Game of Fifteen

`drand48` by first calling `srand48` with an argument of, say, `2`, and then call `drand48` itself three times, `drand48` might instead return `59797`, then `10425`, then `37569`. But if you re-seed `drand48` by calling `srand48` again with an argument of `1`, the next three times you call `drand48`, you'll again get `2728`, then `29785`, then `54710`! See, not so random.

Go ahead and run this program again, this time with a value of, say, `10` for `n`, as in the below; you should see a list of 10 pseudorandom numbers.

```
.....  
./generate 10  
.....
```

Run the program a third time using that same value for `n`; you should see a different list of 10 numbers. Now try running the program with a value for `s` too (e.g., `0`), as in the below.

```
.....  
./generate 10 0  
.....
```

Now run that same command again:

```
.....  
./generate 10 0  
.....
```

Bet you saw the same "random" sequence of ten numbers again? Yup, that's what happens if you don't vary a pseudorandom number generator's initial seed.

Now take a look at `generate.c` itself. (Remember how?) Comments atop that file explain the program's overall functionality. But it looks like we forgot to comment the code itself. Read over the code carefully until you understand each line and then comment our code for us, replacing each `TODO` with a phrase that describes the purpose or functionality of the corresponding line(s) of code. (Know that an `unsigned int` is just an `int` that cannot be negative.) And for more details on `rand` and `srand`, recall that you can execute:

```
.....  
man drand48  
.....
```

and:

```
.....  
man srand48  
.....
```


Problem Set 3: Game of Fifteen

Once done commenting `generate.c`, re-compile the program to be sure you didn't break anything by re-executing the command below.

```
.....  
make generate  
.....
```

If `generate` no longer compiles properly, take a moment to fix what you broke!

Now, recall that `make` automates compilation of your code so that you don't have to execute `clang` manually along with a whole bunch of switches. Notice, in fact, how `make` just executed a pretty long command for you, per the tool's output. However, as your programs grow in size, `make` won't be able to infer from context anymore how to compile your code; you'll need to start telling `make` how to compile your program, particularly when they involve multiple source (i.e., `.c`) files. And so we'll start relying on "Makefiles," configuration files that tell `make` exactly what to do.

How did `make` know how to compile `generate` in this case? It actually used a configuration file that we wrote. Go ahead and look at the file called `Makefile` that's in the same directory as `generate.c`. This `Makefile` is essentially a list of rules that we wrote for you that tells `make` how to build `generate` from `generate.c` for you. The relevant lines appear below.

```
.....  
generate: generate.c  
    clang -ggdb3 -O0 -std=c11 -Wall -Werror -o generate generate.c  
.....
```

The first line tells `make` that the "target" called `generate` should be built by invoking the second line's command. Moreover, that first line tells `make` that `generate` is dependent on `generate.c`, the implication of which is that `make` will only re-build `generate` on subsequent runs if that file was modified since `make` last built `generate`. Neat time-saving trick, eh? In fact, go ahead and execute the command below again, assuming you haven't modified `generate.c`.

```
.....  
make generate  
.....
```

You should be informed that `generate` is already up-to-date. Incidentally, know that the leading whitespace on that second line is not a sequence of spaces but, rather, a tab. Unfortunately, `make` requires that commands be preceded by tabs, so be careful not

Problem Set 3: Game of Fifteen

to change them to spaces, else you may encounter strange errors! The `-Werror` flag, recall, tells `clang` to treat warnings (bad) as though they're errors (worse) so that you're forced (in a good, instructive way!) to fix them.

Now take a look at `find.c`. Notice that this program expects a single command-line argument: a "needle" to search for in a "haystack" of values. Once done looking over the code, go ahead and compile the program by executing the command below.

```
make find
```

Notice, per that command's output, that `make` actually executed the below for you.

```
clang -ggdb3 -O0 -std=c11 -Wall -Werror -o find find.c helpers.c -lcs50 -lm
```

Notice further that you just compiled a program comprising not one but two `.c` files: `helpers.c` and `find.c`. How did `make` know what to do? Well, again, open up `Makefile` to see the man behind the curtain. The relevant lines appear below.

```
find: find.c helpers.c helpers.h
      clang -ggdb3 -O0 -std=c11 -Wall -Werror -o find find.c helpers.c -lcs50 -lm
```

Per the dependencies implied above (after the colon), any changes to `find.c`, `helpers.c`, or `helpers.h` will compel `make` to rebuild `find` the next time it's invoked for this target.

Go ahead and run this program by executing, say, the below.

```
./find 13
```

You'll be prompted to provide some hay (i.e., some integers), one "straw" at a time. As soon as you tire of providing integers, hit ctrl-d to send the program an EOF (end-of-file) character. That character will compel `GetInt` from the CS50 Library to return `INT_MAX`, a constant that, per `find.c`, will compel `find` to stop prompting for hay. The program will then look for that needle in the hay you provided, ultimately reporting whether the former was found in the latter. In short, this program searches an array for some value.

Problem Set 3: Game of Fifteen

At least, it should, but it won't find anything yet! That's where you come in. More on your role in a bit.

It turns out you can automate this process of providing hay, though, by "piping" the output of `generate` into `find` as input. For instance, the command below passes 1,000 pseudorandom numbers to `find`, which then searches those values for `42`.

```
./generate 1000 | ./find 42
```

Note that, when piping output from `generate` into `find` in this manner, you won't actually see `generate`'s numbers, but you will see `find`'s prompts.

Alternatively, you can "redirect" `generate`'s output to a file with a command like the below.

```
./generate 1000 > numbers.txt
```

You can then redirect that file's contents as input to `find` with the command below.

```
./find 42 < numbers.txt
```

Let's finish looking at that `Makefile`. Notice the line below.

```
all: find generate
```

This target implies that you can build both `generate` and `find` simply by executing the below.

```
make all
```

Even better, the below is equivalent (because `make` builds a `Makefile`'s first target by default).

```
make
```

Problem Set 3: Game of Fifteen

If only you could whittle this whole problem set down to a single command! Finally, notice these last lines in `Makefile`:

```
clean:
    rm -f *.o a.out core find generate
```

This target allows you to delete all files ending in `.o` or called `core` (more on that soon!), `find`, or `generate` simply by executing the command below.

```
make clean
```

Be careful not to add, say, `*.c` to that last line in `Makefile`! (Why?) Any line, incidentally, that begins with `#` is just a comment.

search

And now the fun begins! Notice that `find.c` calls `search`, a function declared in `helpers.h`. Unfortunately, we forgot to implement that function fully in `helpers.c`! (To be sure, we could have put the contents of `helpers.h` and `helpers.c` in `find.c` itself. But it's sometimes better to organize programs into multiple files, especially when some functions are essentially utility functions that might later prove useful to other programs as well, much like those in the CS50 Library.) Take a peek at `helpers.c` with, and you'll see that `search` always returns `false`, whether or not `value` is in `values`. Re-write `search` in such a way that it uses linear search, returning `true` if `value` is in `values` and `false` if `value` is not in `values`. Take care to return `false` right away if `n` isn't even positive.

When ready to check the correctness of your program, try running the command below.

```
./generate 1000 50 | ./find 127
```

Because one of the numbers outputted by `generate`, when seeded with `50`, is `127`, your code should find that "needle"! By contrast, try running the command below as well.

```
./generate 1000 50 | ./find 128
```

Problem Set 3: Game of Fifteen

Because `128` is not among the numbers outputted by `generate`, when seeded with `50`, your code shouldn't find that needle. Best to try some other tests as well, as by running `generate` with some seed, taking a look at its output, then piping that same output to `find`, looking for a "needle" you know to be among the "hay".

Incidentally, note that `main` in `find.c` is written in such a way that `find` returns `0` if the needle is found, else it returns `1`. You can check the so-called "exit code" with which `main` returns by executing

```
.....  
echo $?  
.....
```

after running some other command. For instance, assuming your implementation of `search` is correct, if you run

```
.....  
./generate 1000 50 | ./find 127  
echo $?  
.....
```

you should see `0`, since `127` is, again, among the 1,000 numbers outputted by `generate` when seeded with `50`, and so `search` (written by you) should return `true`, in which case `main` (written by us) should return (i.e., exit with) `0`. By contrast, assuming your implementation of `search` is correct, if you run

```
.....  
./generate 1000 50 | ./find 128  
echo $?  
.....
```

you should see `1`, since `128` is, again, not among the 1,000 numbers outputted by `generate` when seeded with `50`, and so `search` (written by you) should return `false`, in which case `main` (written by us) should return (i.e., exit with) `1`. Make sense?

When ready to check the correctness of your program officially with `check50`, you may execute the below.

```
.....  
check50 2015.fall.hacker3.find helpers.c  
.....
```

Incidentally, be sure not to get into the habit of testing your code with `check50` before testing it yourself. (And definitely don't get into an even worse habit of only testing your

code with `check50` !) Suffice it to say `check50` doesn't exist in the real world, so running your code with your own sample inputs, comparing actual output against expected output, is the best habit to get into sooner rather than later. Truly, don't do yourself a long-term disservice!

Anyhow, if you'd like to play with the staff's own implementation of `find`, you may execute the below.

```
~cs50/hacker3/find
```

Sorting

Alright, linear search is pretty meh. Recall from Week 0 that we can do better, but first we'd best sort that hay.

sort

Notice that `find.c` calls `sort`, a function declared in `helpers.h`. Unfortunately, we forgot to implement that function fully too in `helpers.c`! Take a peek at `helpers.c`, and you'll see that `sort` returns immediately, even though `find`'s `main` function does pass it an actual array.

Now, recall the syntax for declaring an array. Not only do you specify the array's type, you also specify its size between brackets, just as we do for `haystack` in `find.c`:

```
int haystack[MAX];
```

But when passing an array, you only specify its name, just as we do when passing `haystack` to `sort` in `find.c`:

```
sort(haystack, size);
```

(Why do you think we pass in the size of that array separately?)

When declaring a function that takes a one-dimensional array as an argument, though, you don't need to specify the array's size, just as we don't when declaring `sort` in `helpers.h` (and `helpers.c`):

Problem Set 3: Game of Fifteen

```
void sort(int values[], int n);
```

Go ahead and implement `sort` so that the function actually sorts, from smallest to largest, the array of numbers that it's passed, in such a way that its running time is in $O(n)$, where n is the array's size. Yes, this running time is possible because you may assume that each of the array's numbers will be non-negative and less than `LIMIT`, a constant defined in `generate.c`. Leverage that assumption! However, realize that the array might contain duplicates.

Now, technically, because we've bounded with a constant the amount of hay that `find` will accept (and because the value of `sort`'s second parameter is bounded by an `int`'s finitely many bits), the running time of `sort`, however implemented, is arguably $O(1)$. Even so, for the sake of this asymptotic challenge, think of the size of `sort`'s input as n .

However, take care not to alter our declaration of `sort`. Its prototype must remain:

```
void sort(int values[], int n);
```

As this return type of `void` implies, this function must not return a sorted array; it must instead "destructively" sort the actual array that it's passed by moving around the values therein. As we'll discuss in Week 4, arrays are not passed "by value" but instead "by reference," which means that `sort` will not be passed a copy of an array but, rather, the original array itself.

Although you may not alter our declaration of `sort`, you're welcome to define your own function(s) in `helpers.c` that `sort` itself may then call.

We leave it to you to determine how best to test your implementation of `sort`. But don't forget that `printf` and GDB are your friends. And don't forget that you can generate the same sequence of pseudorandom numbers again and again by explicitly specifying `generate`'s seed. Before you ultimately submit, though, be sure to remove any such calls to `printf`, as we like our programs' outputs just they way they are!

If you'd like to play with the staff's own implementation of `find`, you may execute the below.

```
~cs50/hacker3/find
```

search

Now that `sort` (presumably) works, it's time to improve upon `search`, the other function that lives in `helpers.c`. Recall that your first version implemented linear search. Rip out the lines that you wrote earlier (sniff) and re-implement `search` as binary search, that divide-and-conquer strategy that we employed in Week 0. You are welcome to take an iterative approach (as with a loop) or a recursive approach (wherein a function calls itself). If you pursue the latter, though, know that you may not change our declaration of `search`, but you may write a new, recursive function (that perhaps takes different parameters) that `search` itself calls. When it comes time to submit this problem set, it suffices to submit this new-and-improved version of `search`; you needn't submit your original version that used linear search.

Here's Zamyła again:

<https://www.youtube.com/watch?v=7DSRJj7qfP8>

The Game Begins

And now it's time to play. The Game of Fifteen is a puzzle played on a square, two-dimensional board with numbered tiles that slide. The goal of this puzzle is to arrange the board's tiles from smallest to largest, left to right, top to bottom, with an empty space in board's bottom-right corner, as in the below.



Sliding any tile that borders the board's empty space into that space constitutes a "move." Although the configuration above depicts a game already won, notice how the tile

Problem Set 3: Game of Fifteen

numbered 12 or the tile numbered 15 could be slid into the empty space. Tiles may not be moved diagonally, though, or forcibly removed from the board. Depicted below is just one configuration from which the puzzle is solvable.



| | | | |
|----|----|----|----|
| 15 | 14 | 13 | 12 |
| 11 | 10 | 9 | 8 |
| 7 | 6 | 5 | 4 |
| 3 | 1 | 2 | |

Okay, navigate your way to `~/workspace/hacker3/fifteen`, and take a look at `fifteen.c`. Within this file is an entire framework for the Game of Fifteen.

Implement God Mode for this game.

First implement `init` in such a way that the board is initialized to a pseudorandom but solvable configuration. Then complete the implementation of `draw`, `move`, and `won` so that a human can actually play the game. But embed in the game a cheat, whereby, rather than typing an integer between 1 and $d^2 - 1$, where d is the board's height and width, the human can also type `GOD` to compel "the computer" to take control of the game and solve it (using any strategy, optimal or non-optimal), making, say, only four moves per second so that the human can actually watch. Presumably, you'll need to swap out `GetInt` for something more versatile. It's fine if your implementation of God Mode only works (bearably fast) for $d \leq 4$; you need not worry about testing God Mode for $d > 4$. Oh and you can't implement God Mode by remembering how `init` initialized the board (as by remembering the sequence of moves that got your program to some pseudorandom but solvable state). Nice try.

To test your implementation, you can certainly try playing it yourself, with or without God Mode enabled. (Know that you can quit your program by hitting ctrl-c.) Be sure that you (and we) cannot crash your program, as by providing bogus tile numbers. And know that, much like you automated input into `find`, so can you automate execution of this game via input redirection if you store in some file a winning sequence of moves for some configuration.

Any design decisions not explicitly prescribed herein (e.g., how much space you should leave between numbers when printing the board) are intentionally left to you. Presumably the board, when printed, should look something like the below (albeit pseudorandom), but we leave it to you to implement your own vision.

Problem Set 3: Game of Fifteen

15 14 13 12

11 10 9 8

7 6 5 4

3 1 2 _

Incidentally, recall that the positions of tiles numbered 1 and 2 should only be swapped (as they are in the 4×4 example above) if the board has an odd number of tiles (as does the 4×4 example above). If the board has an even number of tiles, those positions should not be swapped. Consider, for instance, the 3×3 example below:

8 7 6

5 4 3

2 1 _

Feel free to tweak the appropriate argument to `usleep` to speed up animation. In fact, you're welcome to alter the aesthetics of the game. For (optional) fun with "ANSI escape sequences," including color, take a look at our implementation of `clear` and `check` out http://isthe.com/chongo/tech/comp/ansi_escapes.html for more tricks.

You're welcome to write your own functions and even change the prototypes of functions we wrote. But we ask that you not alter the flow of logic in `main` so that we can automate some tests of your program. If in doubt as to whether some design decision of yours might run counter to the staff's wishes, simply contact your teaching fellow.

If you'd like to play with the staff's own implementation of fifteen on `cloud.cs50.net`, including God Mode, you may execute the below.

```
~cs50/hacker3/fifteen
```

Speaking of God Mode, where to begin? Well, first read up on this Game of Fifteen. Wikipedia is probably a good starting point:

<http://en.wikipedia.org/wiki/N-puzzle>

Problem Set 3:
Game of Fifteen

Then dive a bit deeper, perhaps reading up on an algorithm called A*.

http://en.wikipedia.org/wiki/A*_search_algorithm

Consider using "Manhattan distance" (aka "city-block distance") as your implementation's heuristic. If you find that A* takes up too much memory (particularly for $d \geq 4$), though, you might want to take a look at iterative deepening A* (IDA*) instead:

<http://webdocs.cs.ualberta.ca/~tony/RecentPapers/pami94.pdf>

The staff's own implementation, meanwhile, utilizes an algorithm like that in this paper:

<http://larc.unt.edu/ian/pubs/saml.pdf>

You're welcome to expand your search for ideas beyond those in these papers, but take care that your research does not lead you to actual code. Curling up with others' pseudocode is fine, but please click away if you stumble upon actual implementations (whether in C or other languages).

Alright, get to it, implement this game!